

SHREYANSH BHADANI

AI Engineer | Generative AI & LLM Systems | Backend Development

shreyanshbhadani@gmail.com | [linkedin.com/in/shreyansh-bhadani-77bb28329](https://www.linkedin.com/in/shreyansh-bhadani-77bb28329) | github.com/error404en | shreyansh-37.vercel.app

SUMMARY

Final-year B.Tech student (graduating 2027) building LLM systems end to end: RAG pipelines, multi-model fallback chains, guardrails that keep model output safe to use, and a small language model trained from scratch. Comfortable taking a messy real-world workflow, shipping an AI feature into it, and explaining the trade-offs to non-engineers.

TECHNICAL SKILLS

AI / LLM: RAG, LangChain, embeddings, semantic search, Qdrant, BM25, PyTorch, Transformers (DistilBERT), prompt design, output validation and citation checks, LLM fallback chains

Models & APIs: OpenAI, Gemini, Groq (Llama), Whisper speech-to-text

Backend & Data: Python, FastAPI, REST APIs, PostgreSQL, Supabase, MySQL, Alembic migrations, JWT auth, Row-Level Security

Frontend: React, Next.js, TypeScript, D3.js, Tailwind CSS

Quality & Delivery: pytest, GitHub Actions CI, Docker, Git, Vercel

PROJECTS

JobCopilot

github.com/error404en/Job-Copilot

Next.js, FastAPI, Supabase, Gemini, Groq

Result: Live full-stack AI job-search tool with a 6-model LLM fallback chain, tailored resume export and a persistent coach chat.

- Analyzes postings, scores fit against a resume and goals, and prepares tailored applications without auto-submitting; adds 3-tier ATS company research.
- Resume tailoring uses guardrails against invented skills; generated .docx files are streamed from a download endpoint.

LitLens AI

github.com/error404en/LitLensAI

Next.js 15, React 19, LangChain, Qdrant, Supabase

Result: RAG research assistant over 500+ academic PDFs (35K+ embeddings) with sub-50ms cached queries for 100+ concurrent users.

- 30K+ indexed chunks; Clerk JWT auth with Row-Level Security, Inngest background jobs and Arcjet rate limiting.

EchoLM (Story-SLM)

github.com/error404en/SLM-Creative-Story-

PyTorch, FastAPI, DistilBERT, Next.js, D3.js

Result: 29.63M-parameter GPT-style model trained from scratch on TinyStories: 62.63% next-token accuracy, 287 tok/s on a Tesla T4.

- Final-year team project (3 members). I trained the model and built the Next.js and D3.js story studio (atmosphere modes, creativity and length sliders, token-probability view).
- Built the theme analysis engine: zero-shot atmosphere classification with DistilBERT served via FastAPI.

Equipment Maintenance Triage Assistant

github.com/error404en/Equipment-Maintenance-Triage-Assistant

FastAPI, React, TypeScript, PostgreSQL, Groq

Result: Rules-first triage where AI can never lower a safety priority, and double approvals are blocked in the database (HTTP 409).

- Deterministic rules set severity, BM25 retrieves manual sections, and an optional LLM adds advice that is schema-checked and citation-checked; a named reviewer approves or rejects the draft work order.
- PostgreSQL CHECK constraints, Alembic migrations, pytest suite, GitHub Actions CI and a Docker build.

Also built: EchoSight, accessibility assistant with object detection and voice chat (github.com/error404en/echosight2.0); SecureFS, AES-256-GCM encrypted file storage (github.com/error404en/secure_file).

EXPERIENCE

Generative AI Intern, R&D | Ultimate Flexipack Ltd.

Jun 2026 - Jul 2026

- Built a RAG knowledge assistant for packaging documentation: 300+ documents indexed into 20K+ chunks, cutting manual lookup time by 85%.
- Developed 15+ REST APIs with vector search under 150ms and 90%+ answer accuracy in internal testing.

EDUCATION

B.Tech, Electronics & Communication Engineering (Computer Science specialization)

Expected 2027

Jaypee Institute of Information Technology (JIIT), Noida